

Program Analysis

Fuzzing



Yue Duan

Outline

- Fuzzing recap
- Research paper:
 - Coverage-based Greybox Fuzzing as Markov Chain
 - Directed Greybox Fuzzing



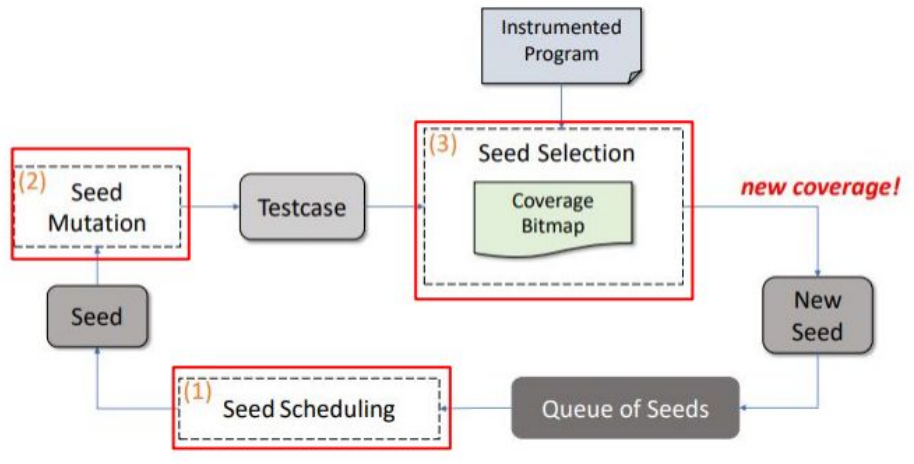
Fuzzing Recap

- Fuzz testing
 - an automated testing technique that uncovers software errors by executing the target program with large number of randomly generated test inputs
 - Three main approaches
 - black-box fuzzing
 - pure random fuzzing
 - grey-box fuzzing
 - leverage some knowledge during testing for guidance
 - white-box fuzzing
 - full knowledge about the target program is needed

Fuzzing Recap

- American Fuzzy Lop (AFL) <https://github.com/google/AFL>
 - coverage-based greybox fuzzer
 - code instrumentation:
 - collect runtime code coverage info
 - can be done at source code and binary level
 - source code: compile with instrumented code
 - binary: QEMU user mode for instrumentation

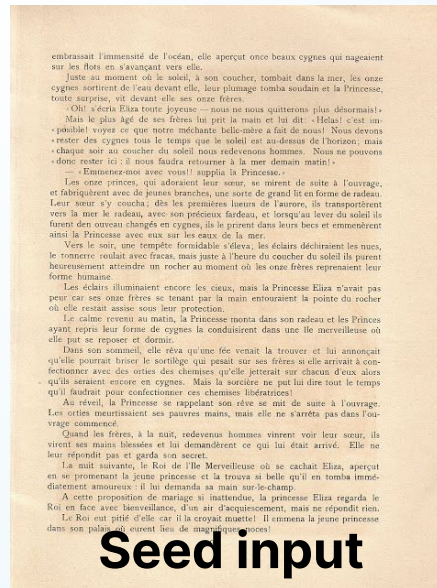
Fuzzing Recap



- Seed scheduling
 - pick the next seed for testing from a set of seed inputs
 - essentially random

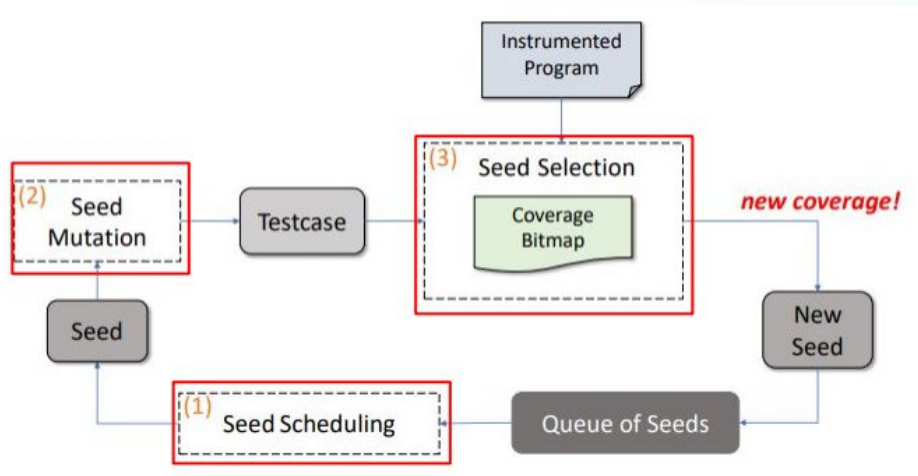
Fuzzing Recap

- Seed Mutation
 - bitflips
 - boundary values (0, -1, 1, INT_MAX, etc)
 - simple arithmetics (add/subtract 1)
 - block deletion
 - block insertion



Fuzzing Recap

- Seed Selection
 - coverage metrics used to define 'interesting' seeds
 - preserve only the interesting inputs for next round
 - coverage bitmap: new basic block discovery



Fuzzing Recap

- Common limitation for grey-box fuzzing:

```
x = int(input())
if x > 10:
    if x < 100:
        print "You win!"
    else:
        print "You lose!"
else:
    print "You lose!"
```

Let's fuzz it!

1 ⇒ "You lose!"
593 ⇒ "You lose!"
183 ⇒ "You lose!"
4 ⇒ "You lose!"
498 ⇒ "You lose!"
48 ⇒ "You win!"

```
x = int(input())
if x > 10:
    if x^2 == 152399025:
        print "You win!"
    else:
        print "You lose!"
else:
    print "You lose!"
```

Let's fuzz it!

1 ⇒ "You lose!"
593 ⇒ "You lose!"
183 ⇒ "You lose!"
4 ⇒ "You lose!"
498 ⇒ "You lose!"
42 ⇒ "You lose!"
3 ⇒ "You lose!"
.....
57 ⇒ "You lose!"

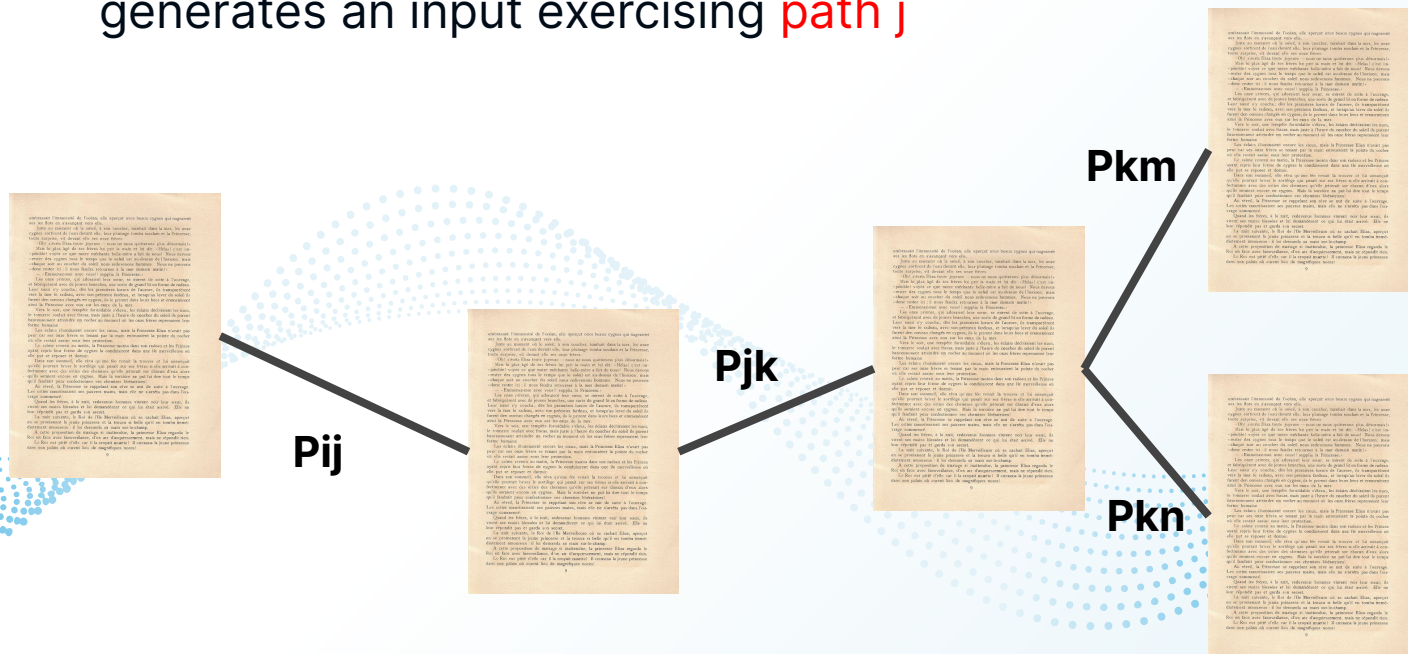
Coverage-based Greybox Fuzzing as Markov Chain

Marcel Böhme , Van Thuan Pham , Abhik Roychoudhury

CCS 2016

Introduction






- Markov chain
 - describes the prob P_{ij} that fuzzing the input exercising **path i** generates an input exercising **path j**




Introduction

- Add **energy** to each state

energy = #fuzz



low energy
(low #fuzz)




high energy
(high #fuzz)

Introduction


How much #fuzz should be generated?

= What is the minimum energy required to expect discovery of new path j from path i?



Challenges


- AFL's power schedule is *constant* in the number of times $s(i)$ for the seed has been chosen for fuzzing.



i

way too much energy


$$P_{ij} = 1/100$$



j


Challenges

- AFL's power schedule is *constant* in the number of times $s(i)$ for the seed has been chosen for fuzzing.




AFLFast

- power schedule is **exponential** in $s(i)$



AFLFast


- power schedule is **exponential** in $s(i)$



Total #fuzz generated: 1

AFLFast


- power schedule is **exponential** in $s(i)$



Total #fuzz generated: 3

AFLFast


- power schedule is **exponential** in $s(i)$



Total #fuzz generated: 7

AFLFast


- power schedule is **exponential** in $s(i)$



Total #fuzz generated: 15

AFLFast


- power schedule is **exponential** in $s(i)$



Total #fuzz generated: 31

AFLFast

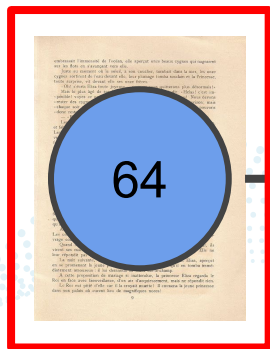
- power schedule is **exponential** in $s(i)$



Total #fuzz generated: 63


AFLFast

- power schedule is **exponential** in $s(i)$



$$P_{ij} = 1/100$$

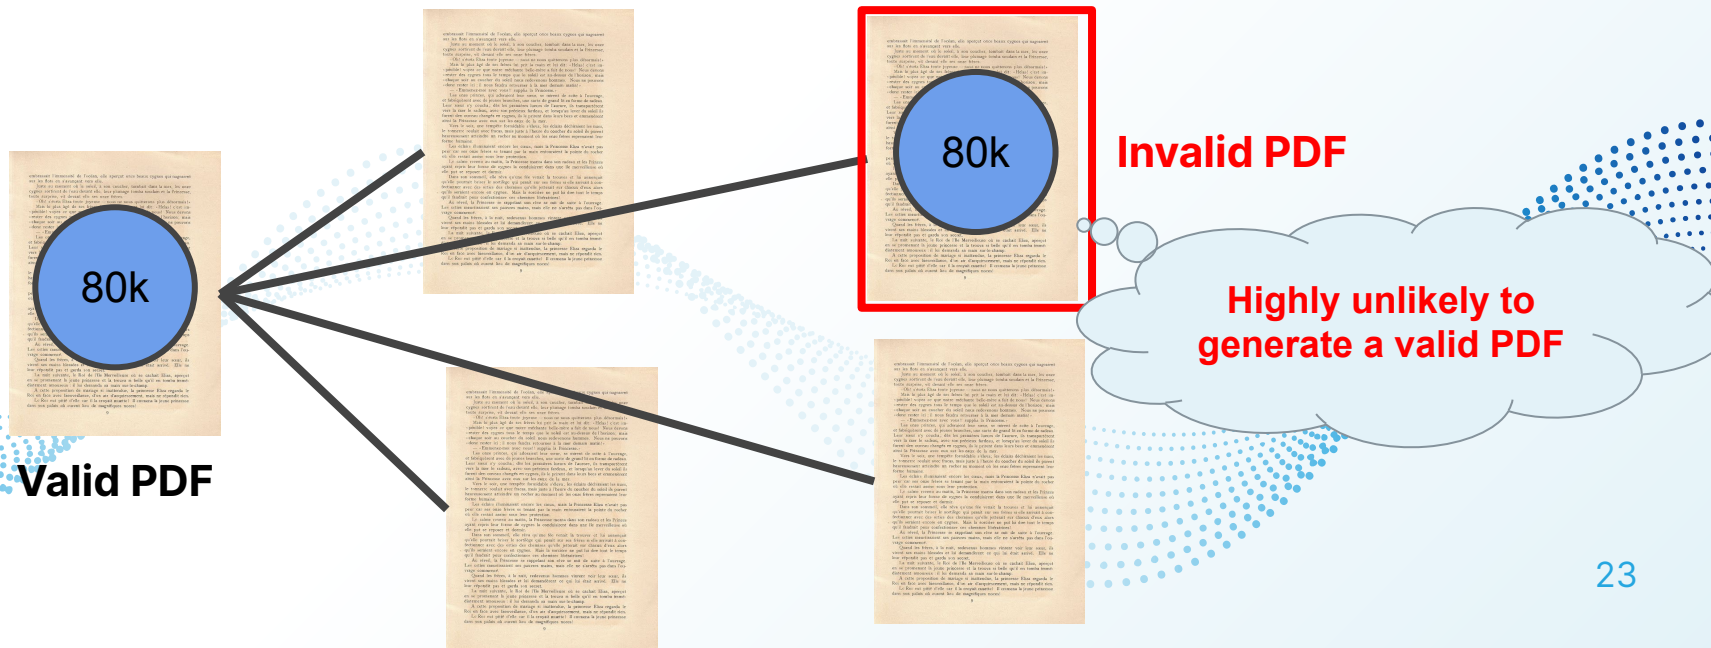
Discovered with much
smaller test cases!



Total #fuzz generated: 127

Challenges

- AFL's power schedule is *constant* in the number of times $s(i)$ for the seed has been chosen for fuzzing.
- AFL's power schedule always assigns *high* energy



Challenges

- AFL's power schedule is *constant* in the number of times $s(i)$ the seed has been chosen for fuzzing.
- AFL's power schedule always assigns *high* energy
 - Too much energy assigned to high-frequency paths

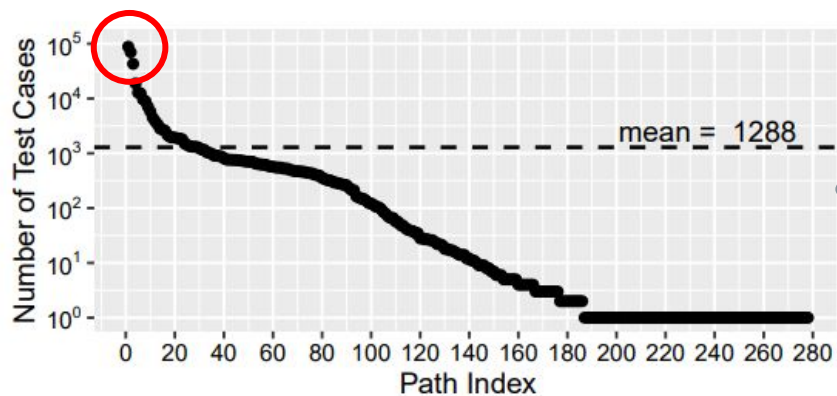


Figure 1: #Fuzz exercising a path (on a log-scale) after running AFL for 10 minutes on the nm-tool.

Most fuzz exercises the same few paths

AFLFast

- AFLFast's power schedule assigns energy that is inversely proportional to the density of the stationary distribution
 - assign low energy to high-frequency paths
 - assign high energy to low-frequency paths
 - approximate the density of distribution for a state i by counting the number of fuzz $f(i)$ that exercises the path i

Power Schedules

- AFL: constant power schedule
 - $p(i) = \mathbf{a}(i)$
 - $\mathbf{a}(i)$ is how AFL judges fuzzing time for path i (~1 min)
- AFLFast:
 - spend *more energy* on low-frequency paths
 - *less energy* on high-frequency paths
 - spend the *minimum energy* required to discover a new state

$$p(i) = \begin{cases} 0 & \text{if } f(i) > \mu \\ \min\left(\frac{\alpha(i)}{\beta} \cdot 2^{s(i)}, M\right) & \text{otherwise.} \end{cases}$$

$$p(i) = \min\left(\frac{\alpha(i)}{\beta} \cdot \frac{2^{s(i)}}{f(i)}, M\right)$$

Search Strategies

- AFL
 - chooses the seeds in the order they are added
 - after the last seed, begin with the first
- AFLFast
 - prioritizes seeds that
 - exercise low-frequency paths
 - have been chosen less often
 - chooses each seed at most once per cycle

Evaluation

- Binutils

Table 1: CVE-IDs and Exploitation Type

Vulnerability	Type
CVE-2016-2226	Exploitable Buffer Overflow
CVE-2016-4487	Invalid Write due to a Use-After-Free
CVE-2016-4488	Invalid Write due to a Use-After-Free
CVE-2016-4489	Invalid Write due to Integer Overflow
CVE-2016-4490	Write Access Violation
CVE-2016-4491	Various Stack Corruptions
CVE-2016-4492	Write Access Violation
CVE-2016-4493	Write Access Violation
CVE-2016-6131	Stack Corruption
Bug 1	Buffer Overflow (Invalid Read)
Bug 2	Buffer Overflow (Invalid Read)
Bug 3	Buffer Overflow (Invalid Read)

Evaluation

- General results

Vulnerability	AFL	<i>AFL-Fast</i>	Factor
CVE-2016-2226	> 24.00 h	3.85 h	N/A
CVE-2016-4487	2.63 h	0.46 h	5.8
CVE-2016-4488	6.92 h	0.98 h	7.0
CVE-2016-4489	10.68 h	2.78 h	3.8
CVE-2016-4490	3.68 h	0.41 h	9.1
CVE-2016-4491	> 24.00 h	4.74 h	N/A
CVE-2016-4492	12.18 h	0.87 h	14.1
CVE-2016-4493	4.48 h	1.00 h	4.5
CVE-2016-6131	> 24.00 h	5.48 h	N/A
Bug 1	20.43 h	3.38 h	6.0
Bug 2	20.91 h	2.89 h	7.2
Bug 3	> 24.00 h	5.07 h	N/A

Figure 8: Time to expose the vulnerability.

Evaluation

- Power schedules

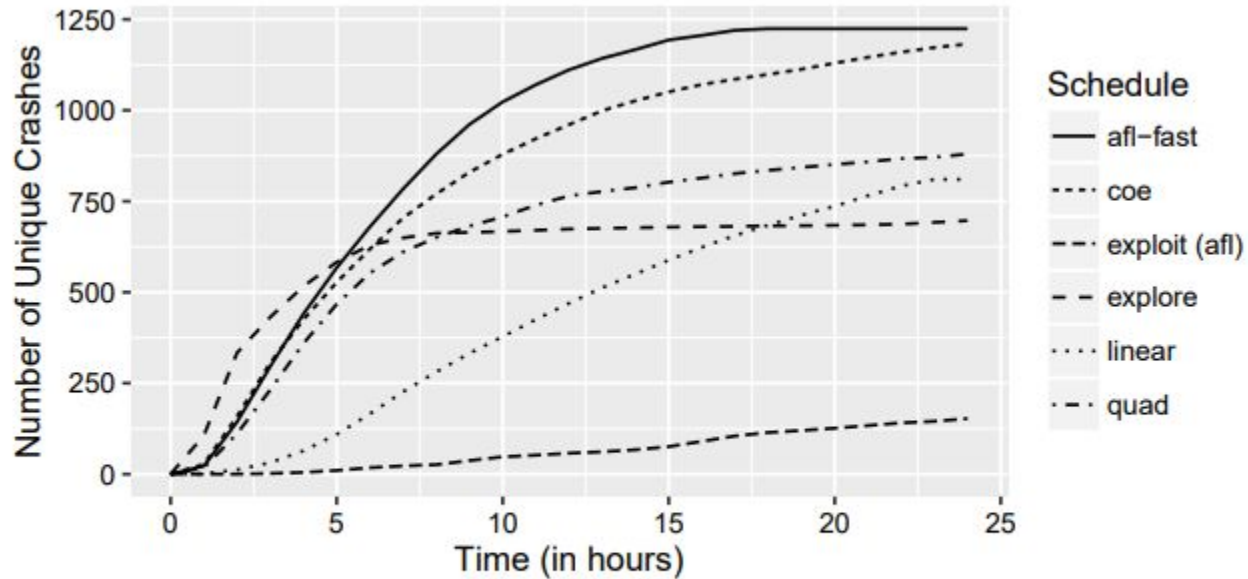


Figure 10: #Crashes over Time (Schedules).

Directed Greybox Fuzzing

Marcel Böhme, Van Thuan Pham, M.D Nguyen, Abhik Roychoudhury

CCS 2017

Motivation

- Grey-box fuzzing is frequently used
 - state-of-the-art in automated vulnerability detection
 - extremely efficient
 - all program analysis before/at instrumentation time
 - start with a seed, choose a seed file, fuzz it
 - add to corpus only if new input increases coverage
- **Cannot be directed!**

Motivation

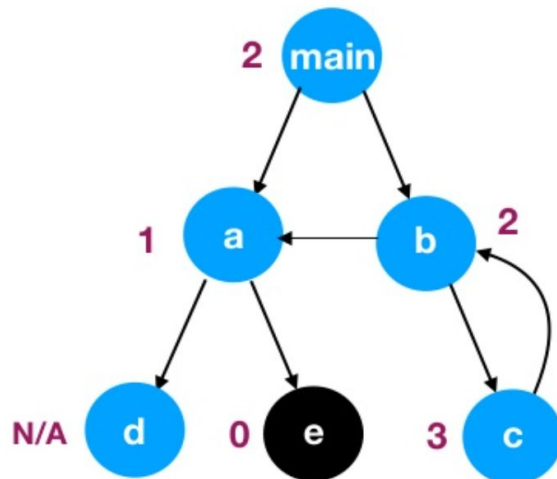
- Directed fuzzing
 - patch testing: reach changed statements
 - crash reproduction: exercise stack traces
 - SA report verification: reach 'dangerous' locations
 - etc

Overview of AFLGo

- Directed fuzzing as **optimization problem**
 - instrumentation time:
 - extract call graph and control-flow graphs
 - for each BB, compute distance to target locations
 - instrument program to aggregate distance values
 - Runtime
 - collect coverage and distance information
 - decide **how long to be fuzzed** based on distance
 - if input is closer to the targets, fuzz longer
 - if input is further away, fuzz shorter

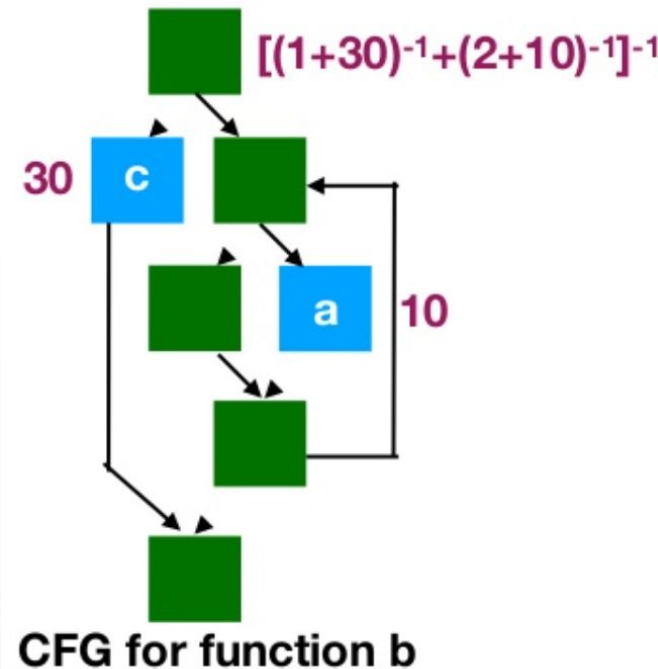
Instrumentation

- Function-level target distance
 - use call graph
 - identify target functions in CG
 - compute the harmonic mean of the lengths of the shortest path to targets



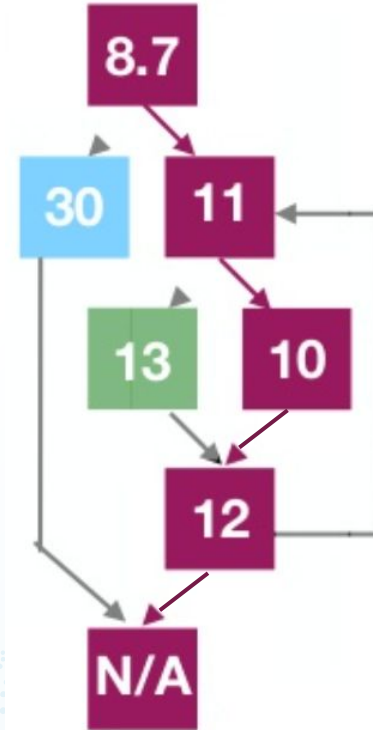
Instrumentation

- BB-level target distance
 - use CFG
 - identify target BBs and assign **distance 0**
 - identify BBs that call functions and assign **10*FLTD**
 - for each BB, compute harmonic mean of
 - **(length of shortest path to any function-calling BB + 10*FLTD)**



Runtime

- Seed distance
 - from instrumented binary
 - two 64-bit shared memory entries
 - aggregated BB-level distance values
 - number of executed BBs



Seed distance: $10.4 = (8.7 + 11 + 10 + 12)/4$

Directed Fuzzing

- Assign **more energy** to seeds that are closer to the given targets
- Problem
 - if always assign more energy to closer seeds
 - likely reach only a local minimum distance but never a global minimum distance
- Solution (simulated annealing)
 - sometimes assign more energy to further-away seeds
 - approaches global minimum distance

Evaluation

- Patch testing: reach changed statements
 - state-of-the-art
 - KATCH (based on symbolic execution)
 - Experiment setup
 - reuse KATCH benchmark
 - measure path coverage
 - measure vulnerability detection

Evaluation

- Patch testing: reach changed statements

	#Changed Basic Blocks	#Uncovered Changed BBs	KATCH	AFLGo
<i>Binutils</i>	852	702	135	159
<i>Diffutils</i>	166	108	63	64
Sum	1018	810	198	223

	KATCH #Reports	#Reports ¹⁴	AFLGo #New Reports	#CVEs
<i>Binutils</i>	7	4	12	7
<i>Diffutils</i>	0	N/A	1	0
Sum	7	4	13	7

Evaluation

- Crash reproduction: exercise stack trace
 - state-of-the-art:
 - BugRedux (based on symbolic execution)
 - Experiment setup
 - reuse original BugRedux benchmark
 - determine whether or not crash can be reproduced

Subjects	BugRedux	AFLGo	Comments
sed.fault1	✗	✗	Takes two files as input
sed.fault2	✗	✓	
grep	✗	✓	
gzip.fault1	✗	✓	
gzip.fault2	✗	✓	
ncompress	✓	✓	
polymorph	✓	✓	

Thank you!

Questions?