

---

---

# Cybersecurity in Autonomous Driving

---

---

Yue Duan

# Too Afraid to Drive: Systematic Discovery of Semantic DoS Vulnerability in Autonomous Driving Planning under Physical-World Attacks

Ziwen Wan, Junjie Shen, Jalen Chuang, Xin Xia,  
Joshua Garcia, Jiaqi Ma, and Qi Alfred Chen

*AS<sup>2</sup>Guard*


Autonomous & Smart Systems  
Guard Research Group

UCI

UCLA

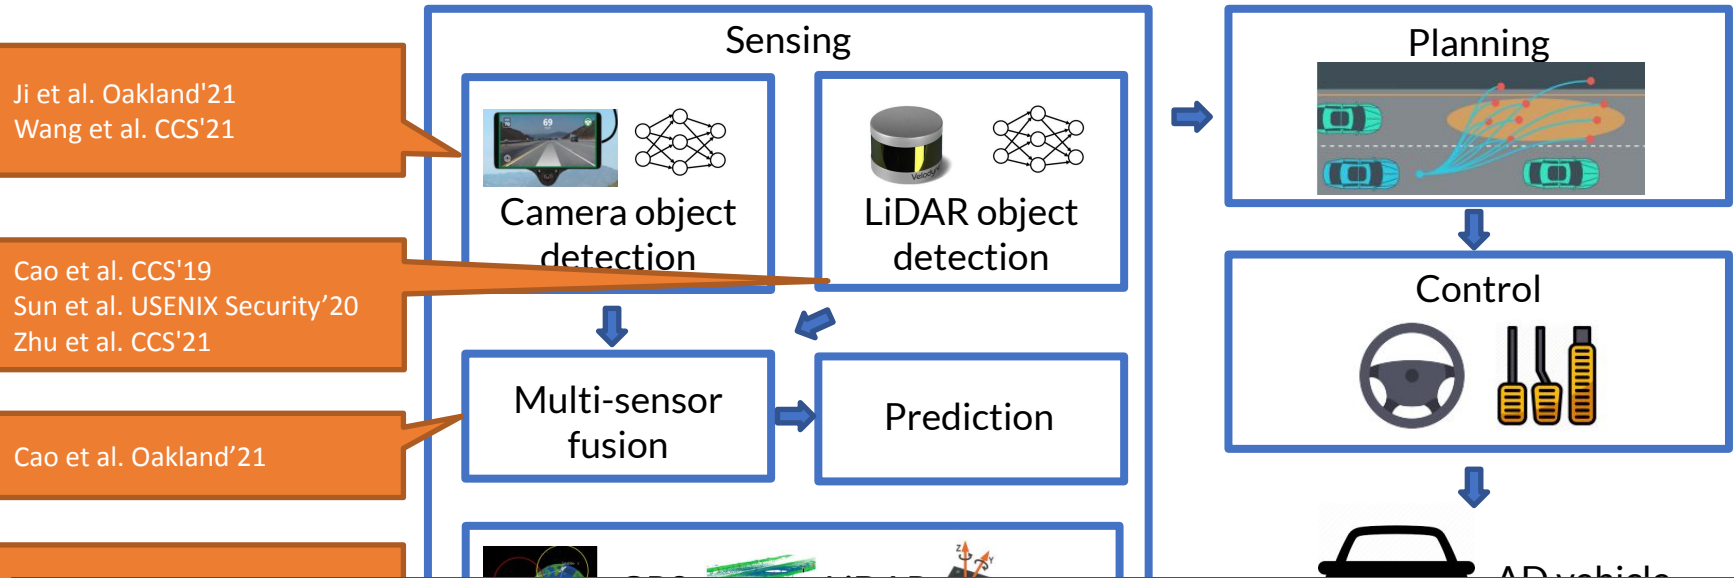
# The rise of Autonomous Driving (AD) vehicles

- **High-level** autonomous driving vehicles are already providing services **without safety drivers**.



# Current status of AD security research

- We have witnessed security problems in high-level AD systems.



**Question:** Could planning (critical driving decision-making) also be vulnerable and thus exploitable to external attackers?

# Our focus: Semantic vulnerability in AD planning

- **Definition:** causing planning to change a normal driving decision to an unexpected one




Attacker

Manipulate external  
AD system inputs



AD system

Sensing



Perception &  
prediction



Localization

Planning



Bugs,  
design  
flaws

Our focus in this work.

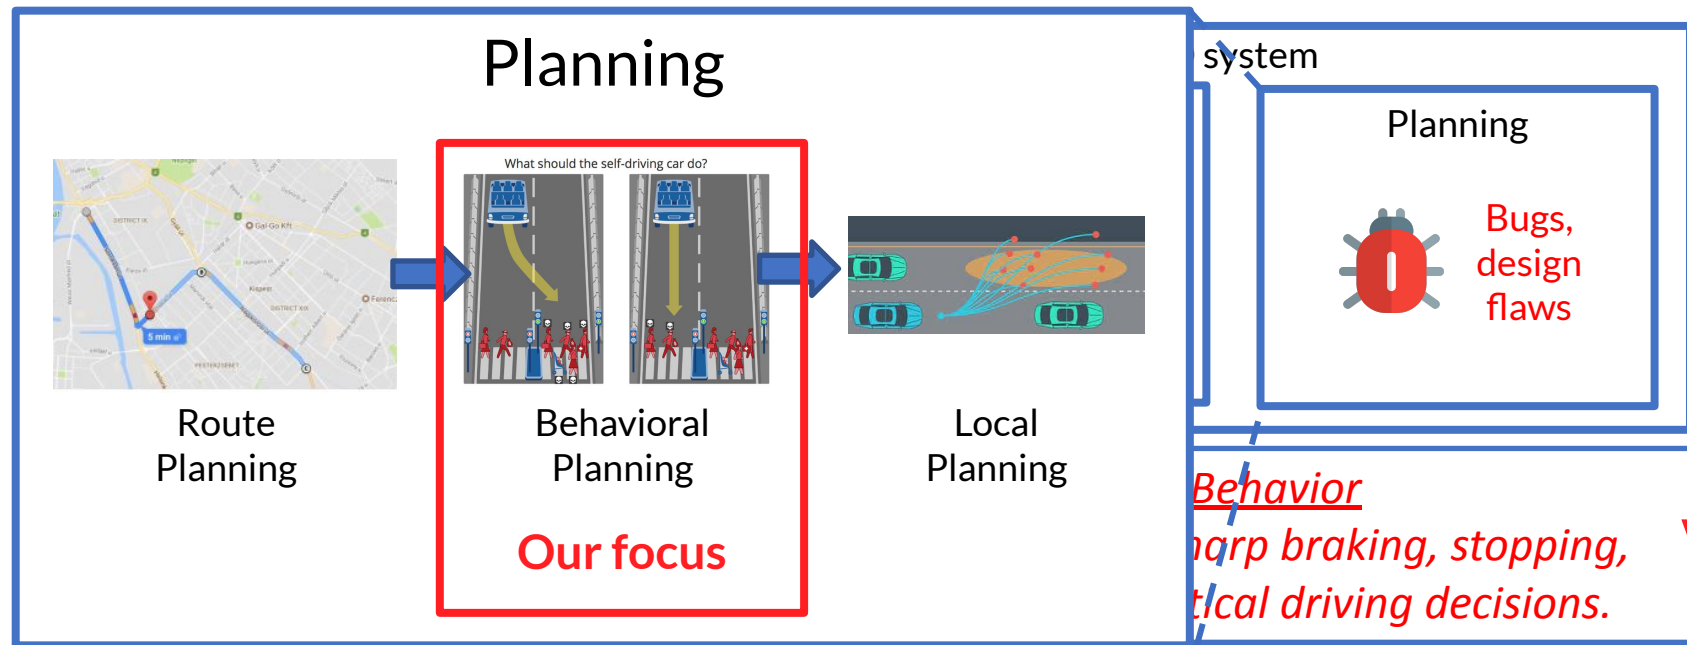
Also referred to as  
semantic DoS vulnerability

Overly-Conservative Behavior

- e.g., unnecessary sharp braking, stopping,  
giving up mission-critical driving decisions.

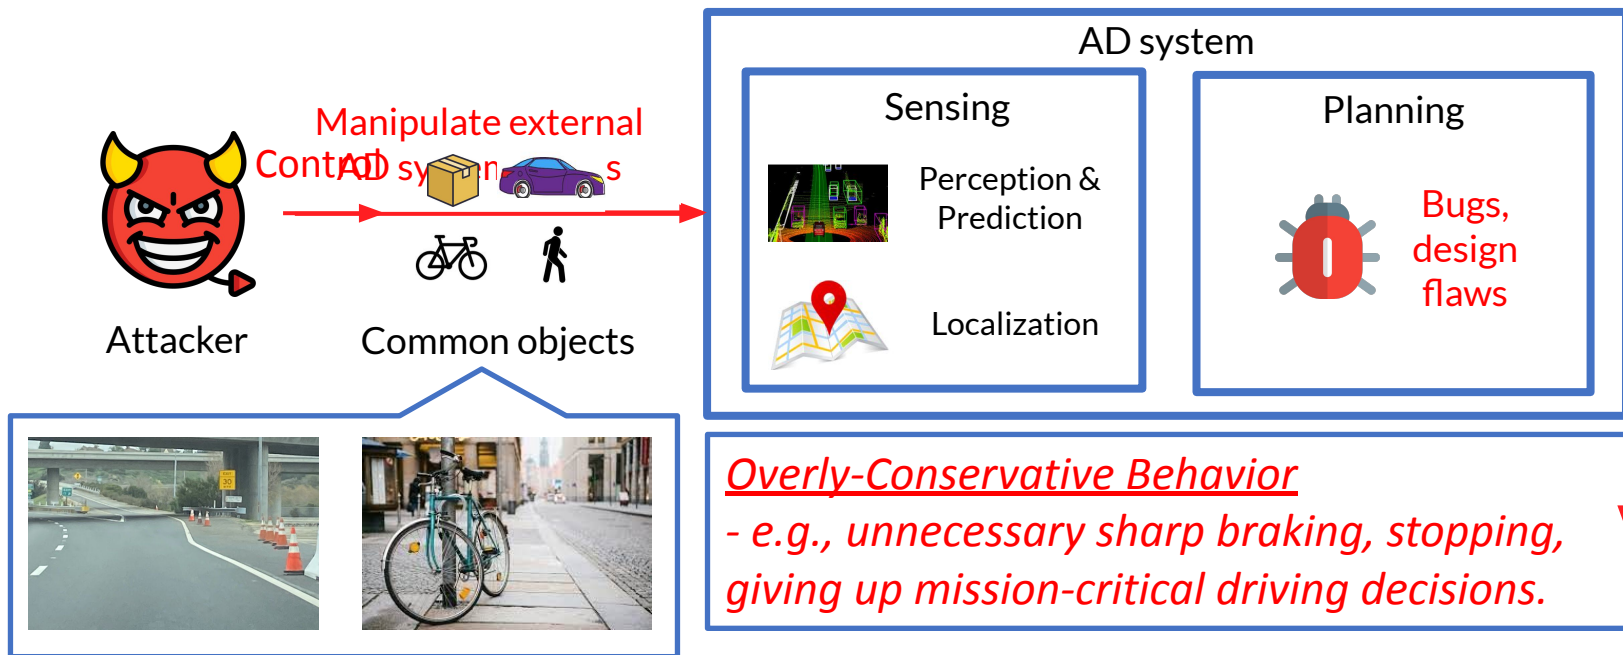
# Our target: Behavioral Planning (BP)

- Functionality of BP: Makes mission-critical driving decisions, e.g., collision avoidance, lane changing



# Threat model

- **Attack vector:** attacker-controllable common roadside objects
  - e.g., dumped cardboard boxes, parked bikes on the road side



# Consequence of semantic DoS vulnerability

## Consequences



Bad user  
experience



Safety



Block traffic



Law violation in  
specific places

*Overly-Conservative Behavior*

*- e.g., unnecessary sharp braking, stopping,  
giving up mission-critical driving decisions.*




# Semantic DoS vulnerability demo



As a human driver, how should you react to this scenario at the highway off-ramp?

- ☐ Ignore them?
- ☐ Slightly slow down?




**Now let's look into a demo we created with Autoware.AI.**



Two pictures around our campus.

# Attack scenario setup

# Root cause of the DoS vulnerability




Drivable area (minimal value is  $(3.5 - 2 \times 1.2) < \text{car width (2.11m)}$ )  
The AD vehicle thinks there is not enough space

# PlanFuzz: Design challenges

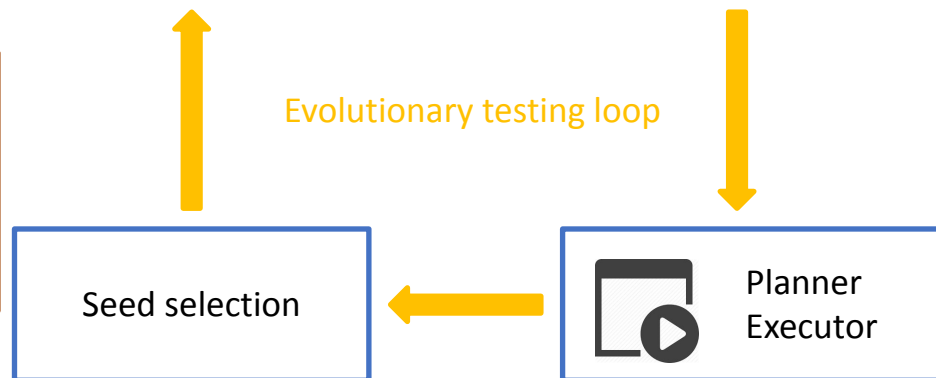
- We design *PlanFuzz*, a novel dynamic testing tool to automate the semantic DoS vulnerabilities discovery

**Challenge 2:** How to generate inputs that satisfy domain constraints?




**Challenge 1:** How to judge a driving decision is *overly-conservative*?

**Challenge 3:** How to design feedback to efficiently guide the testing ?



# Solution: Planning Invariant (PI)

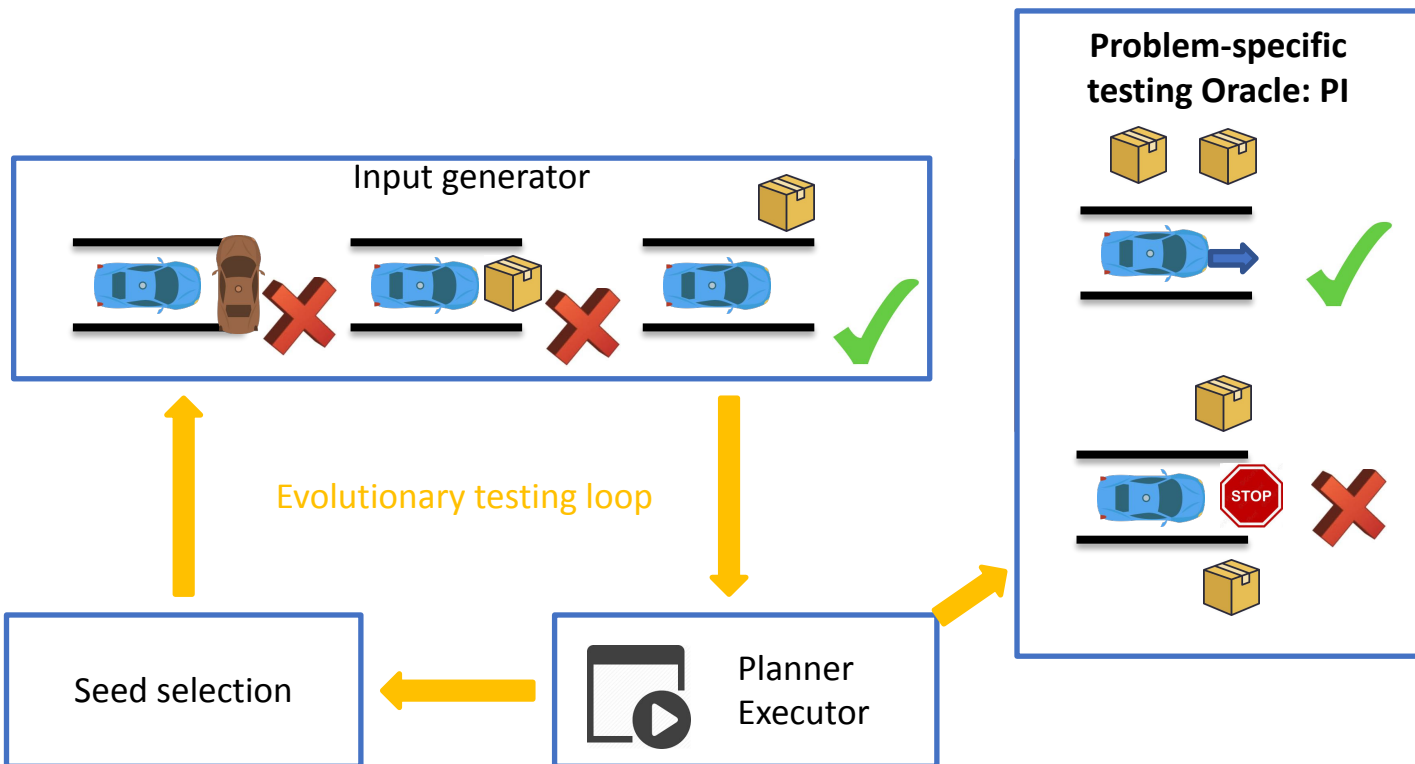
- To address challenge 1 (lack of testing oracles for semantic DoS vuln), we design planning invariant
  - Planning Invariants (PI) = planning scenario + desired planning behavior + attacker-controllable changes
  - **Systematically** define PIs under 8 diverse scenarios with temporal logic to constraint static objects, and **moving** pedestrian/vehicles



# Solution: Planning Invariant (PI)

**Challenge 2:** How to generate inputs that satisfy domain constraints?

**Challenge 3:** How to design feedback to efficiently guide the testing?




# Solution: PI-aware physical-object generation


## Input generation:

- Satisfy domain-specific constraints
- Maintain diversity and inheritance during mutation

## PI-aware physical-object generation



## Problem-specific testing Oracle: PI



**Challenge 3:** How to design feedback to efficiently guide the testing?

Evolutionary testing loop

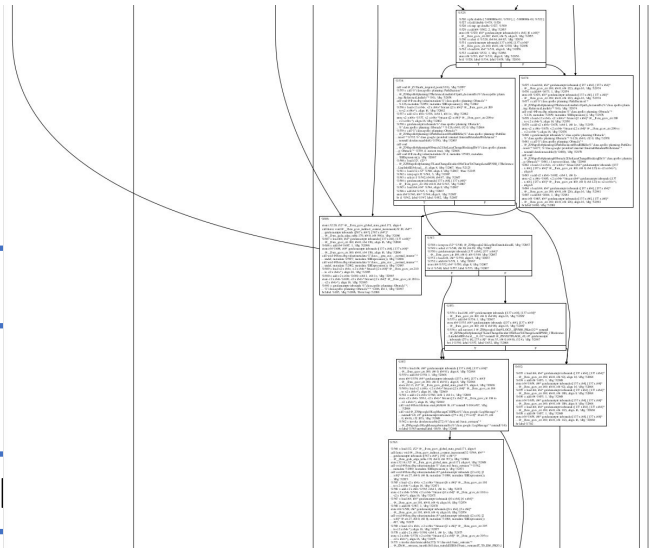
Seed selection



Planner  
Executor

# Solution: BP vulnerability distance

- To address challenge 3 (lack of efficient guidance)
  - We propose **BP vulnerability distance**, which is a **gray-box** guidance.



Tiny fraction of Apollo lane changing control flow graph

**Key idea:** Use the distance between operands in decision-related predicates to guide driving decision changes

## Offline static analysis:

- Extract control/data dependency
- Generate BP vuln. distance profile for instrumentation

## Online dynamic analysis:

- Calculate BP vuln. dist. at runtime



# Solution: BP vulnerability distance

Offline analysis & instrument phase



Source code

Control & data  
dependency  
analyzer

BP vuln. Distance  
profile generator


Challenge 2: How to  
design feedback to

BP vuln. Trace  
instrumentor


Online vulnerability  
testing phase

PI-aware physical-object generation


Static property  
generation



PI-constraint  
enforcement



Dynamic property  
generation




Evolutionary testing loop

Seed selection based  
on BP vulnerability distance




Instrumented  
planner  
executor

Problem-specific  
testing Oracle: PI




# Evaluations: DoS semantic vulnerabilities discovery


- **9 previously unknown** semantic DoS vulnerabilities from **3 BP implementations** of Baidu Apollo and Autoware.AI (full-stack open-source AD software)
  - Causes: 1 due to implementation bug, 8 due to overly-conservative planning parameters (e.g., safety buffer, angle threshold) & overly-conservative estimation of surrounding object intentions (e.g., from pedestrians, parked bicycles)
- **Diverse** driving scenarios
  - 28,789 BP decision snapshots from 40 driving traces & 8 different scenario types



Lane following



Lane changing





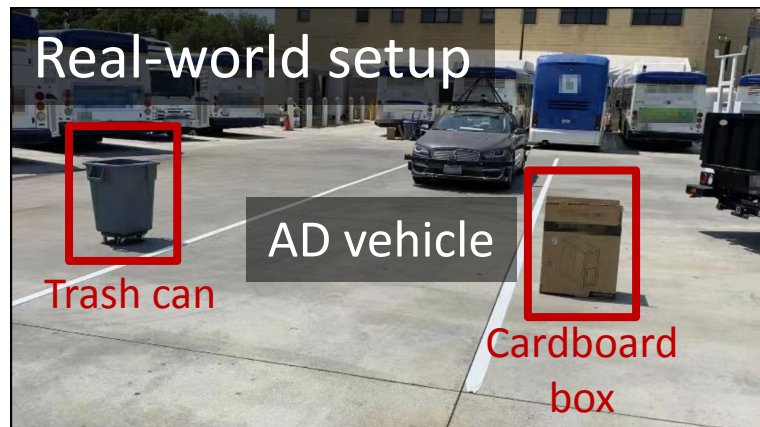
Lane borrowing



Intersection passing

More evaluations in the paper...

# Exploitation case studies



# Conclusion

**First** to perform AD planning-specific semantic vulnerability discovery with a **domain-specific vulnerability definition** and a **practical threat model**

- Design *PlanFuzz*, a **novel dynamic testing** approach that addresses various problem-specific design challenges
- We evaluate *PlanFuzz* on **two** practical open-source **full-stack** AD systems and discover **9** previously-unknown DoS vulnerabilities
- Perform exploitation case studies of **3 diverse driving scenarios** with simulation and driving traces collected from **a real AD vehicle**
- Inform **24 companies** developing AD vehicles

**Thank you!**  
**Question?**